This is the accepted manuscript of:

Benefits in relaxing the power capping constraint

Cesarini, Daniele; Bartolini, Andrea; Benini, Luca

Source: ACM International Conference Proceeding Series, v Part F132205, September 9, 2017, 1st Workshop on AutotuniNg and aDaptivity AppRoaches for Energy efficient HPC Systems, ANDARE 2017 - A Workshop part of PACT 2017

©ACM 2017. This version of the work is posted here for your personal use. Not for redistribution. The definitive Version of Record was published in https://doi.org/10.1145/3152821.3152878

Benefits in Relaxing the Power Capping Constraint

Daniele Cesarini DEI, University of Bologna Bologna, Italy daniele.cesarini@unibo.it Andrea Bartolini DEI, University of Bologna Bologna, Italy a.bartolini@unibo.it Luca Benini
IIS, Swiss Federal Institute of
Technology
Zürich, Switzerland
lbenini@iis.ee.ethz.ch

ABSTRACT

In this manuscript we evaluate the impact of HW power capping mechanisms on a real scientific application composed by parallel execution. By comparing HW capping mechanism against static frequency allocation schemes we show that a speed up can be achieved if the power constraint is enforced in average, during the application run, instead of on short time periods. RAPL, which enforces the power constraint on a few ms time scale, fails on sharing power budget between more demanding and less demanding application phases.

KEYWORDS

power capping, DVFS, HPC, RAPL, monitoring, P-states, power management, hardware performance counters

ACM Reference Format:

Daniele Cesarini, Andrea Bartolini, and Luca Benini. 2017. Benefits in Relaxing the Power Capping Constraint. In ANDARE '17: 1st Workshop on AutotuniNg and aDaptivity AppRoaches for Energy efficient HPC Systems, September 9, 2017, Portland, OR, USA. ACM, New York, NY, USA, 6 pages. https://doi.org/10.1145/3152821.3152878

1 INTRODUCTION

The pace dictated by the Moore's law on technological scaling has provided a constant increase in the performance of computing devices. However with the breakdown of Dennard scaling, this has come with an increase in the power consumption.

Supercomputers are the cutting edge of computing performance. Supercomputers integrate hundred thousand of the most powerful processors and accelerators and they are periodically ranked based on their peak performance [1]. Until June 2016, every new most powerful supercomputer in the world (1st in the Top500 list) has marked an increase in its power consumption. Reaching in 2013, with Tianhe-2, 17.8 MWatts of IT peak power consumption, which increases to 24MWatt when considering also the cooling power [8]. This has set the record for the power consumption of a single supercomputer installation, reaching the practical limit in power provisioning which is 20MWatt [3]. The today's most powerful supercomputer (Taihulight) consumes 15.4 MWatt underlining the fact that performance increase is nowadays possible only at a fix power budget: as a matter of fact supercomputers are power limited!

To ensure this power budget during design time, it requires considering the worst-case power consumption of the computing resources. However, supercomputer workloads rarely cause worstcase power consumption during their life making worst-case design approaches a bad choice which decreases the average supercomputing performance. Power capping approaches support the design for "the average power consumption case" by packing more computing capacity, with a feasible power budget under average workload and dynamically reducing the performance during the execution of workloads with peak power consumption.

At the basis of these approaches there is the capability of computing elements to trade off performance with power consumption. Dynamic and Voltage Frequency states (DVFS) (ACPI P-states [13]) allow the reduction of the power consumption during active computation. Dynamic power management policies take advantage of these HW states to create feedback loops adapting the performance to workload phases aiming to reduce the energy consumption or ensure a specific power and thermal budget. Pure software implementations of these policies have clear software advantages but need to be executed on the same computational resources, interrupting the application flow and causing overheads.

Recently vendors have added HW components to implement in HW these policies allowing a more fine-grain control. Intel Running Average Power Limits (RAPL) technology can enforce in hardware a given power limit. This is done exploiting power sensors and power management knobs. Current RAPL implementations aim to enforce a constant power consumption within a ten-milliseconds time window. However this is far below the timescale of supercomputing centres where the power budget must be respected in large time windows coming from the power grid and supplier [6, 7]. Differently, RAPL enforces the power cap for every application phase without taking into account the real power efficiency of these phases.

In this exploration work, we compare RAPL mechanism against fixed frequency assignment (which causes the same power consumption as the one of the RAPL case, but in average for the entire application run).

In detail, we analyze the efficacy (measured as application time-to-solution) of power capping to maintain the same power budget using (i) RAPL controller and (ii) a simple fixed frequency allocation that statically assigns core's frequencies for the entire application time. In the first case RAPL imposes the power constraint every ten milliseconds modulating the operating point (i.e. clock frequency) to workload changes with a similar time granularity. The second setup on the contrary, uses fixed frequency for the entire application time chosen to obtain in average for the entire duration of the application the same power consumption as RAPL.

The paper is organized as follows. Section 2, presents related works, Section 3 characterizes HPC architectures and their power manager used to power constraint the system. Section 4 defines our architecture and application targets presenting the monitoring infrastructure. While section 5 shows the methodology used for the exploration and our experimental results.

2 RELATED WORK

Several approaches in the literature have proposed mechanisms to constrain the power consumption of large-scale computing infrastructures. These can be classified into two main families. Approaches in the first class use predictive models to estimate the power consumed by a job before its execution. At job scheduling time, this information is used to allow into the system jobs that

satisfy the total power consumption budget. Hardware power capping mechanisms like RAPL are used to ensure that the predicted budget is respected during all the application phases and to tolerate prediction errors in the job average power consumption estimation [4, 5, 19]. Approaches in the second class distribute a slice of the total system power budget to each active computing element. The per-compute element power budget is ensured to hardware power capping mechanisms like RAPL. The allocation of the power consumption budget to each compute node can be done statically or dynamically [9, 10, 14, 18]. It is the goal of the run-time to trade off power reduction with application performance loss. The GEOPM [9] runtime developed by Intel is an open source, plugin extensible runtime for power management. GEOPM implement a plugin for power balancing to improve performance in power constraint systems reallocating power on sockets involved in the critical path of the application. Authors in [20] quantitatively evaluated RAPL as a control system in term of stability, accuracy, settling time, overshoot, and efficiency. In this work, authors evaluate only the proprieties of RAPL mechanism without considering other power capping strategies and how can vary application workload.

3 HPC ARCHITECTURES AND POWER MANAGEMENT SYSTEMS

3.1 HPC Architectures

HPC systems are composed of tens to thousands computational nodes interconnected with a low-latency high-bandwidth network. Nodes are usually organized in sub-clusters allocated at execution time from the system scheduler according to the user request. Subclusters have a limited lifetime, after which resources are released to the system scheduler. Users request resources through a batch queue system, where they submit applications to be executed. Even a single node can be split in multiple resources shared among users. The single indivisible units in a HPC machine are CPUs, memory and possibly accelerators (GPGPU, FPGA, Many-core accelerator, etc.).

HPC applications typically use the Single Program Multiple Data (SPMD) execution model, where the same application executable is instanced multiple time on different nodes of the cluster; each instance works on a partition of the global workload and communicates with other instances to orchestrate subsequent computational steps. For this reason, a HPC application can be seen as the composition of several tasks executed in a distributed environment which exchanges messages among all the instances. Achieving high-performance communication on distributed applications in large clusters is not an easy task. The Message-Passing Interface (MPI) runtime responds to these demands by abstracting the level of network infrastructure using a simple but high-performance interface for communication that can scale up on thousands of nodes.

HPC machines are extreme energy consumers and server rooms require a proportioned cooling system to avoid overheating situations. The extreme working conditions of this kind of machines bring a lot of inefficiencies in terms of energy and thermal control, that turn in computational performance degradation. Hardware power managers are becoming a fundamental component to control power utilization using different strategies in order to reduce energy waste and, at the same time, assure a safe thermal environment.

3.2 Power Management in HPC Systems

Nowadays, operating systems can communicate with different hardware power managers through an open standard interface called Advanced Configuration and Power Interface (ACPI) [13]. In this work, we focus on ACPI implementation of Intel architecture, since most HPC machines (more than 86% in [1]) are based on Intel CPUs.

Intel implements the ACPI specification defining different component states which a CPU can use to reduce power consumption. Today's CPU architectures are composed of multiple processing elements (PE) which communicate through a network subsystem that interconnect PEs, Last Level Cache (LLC), Integrated Memory Controllers (IMC) and other uncore components. Intel architecture optimizes ACPI using different power saving levels for cores and uncore components. The ACPI standard defines P-states to select DVFS operating points targeting the reduction of active power, while also defines specific idle states to manage power consumption during inactivity time of the CPU. In our work, we consider only P-states to manage DVFS control knob, this because HPC applications do not manifest idle time during the execution.

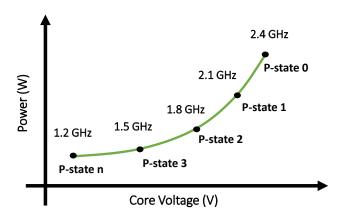


Figure 1: DVFS levels and Intel P-states

Intel P-States show in figure 1, defining a number of levels which are numbered from 0 to n where n is the lowest frequency and 0 is the highest frequency with the possibility to take advantage of Turbo Boost technology. Turbo Boost is an Intel technology that enables processors to increase their frequency beyond the nominal via dynamic control of clock rate. The maximum turbo frequency is limited by the power consumption, thermal limits and the number of cores that are currently using turbo frequency. Since Haswell, Intel cores allow independent per-core P-state.

3.2.1 Linux Power Management Driver. Intel P-states are managed by a power governor implemented as a Linux kernel driver. By default on Linux system, Intel architectures are managed by a kernel module called <code>intel_pstate</code> developed by Intel. But <code>intel_pstate</code> driver does not support a governor that allows users to select percore fixed frequency. Differently, the standard power management driver of Linux <code>acpi-cpufreq</code> does it. <code>acpi-cpufreq</code> is similar to Intel driver but implement a large set of governors. The available governors are:

- (1) *powersave*: this governor runs the CPU always at the minimum frequency.
- (2) performance: runs the CPU always at the maximum frequency.
- (3) userspace: runs the CPU at user specified frequencies.
- (4) ondemand: scales the frequency dynamically according to current load [17].
- (5) *conservative*: similar to *ondemand* but scales the frequency more gradually.

In our work we use *acpi-cpufreq* driver with *userspace* governor that allows us to select per-core fixed frequency.

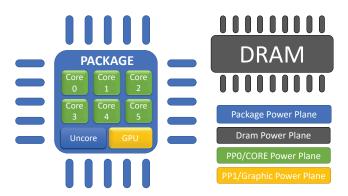


Figure 2: Intel RAPL design with the identification of power domains

3.3 Hardware Power Controller

Today's CPU architectures implement reactive hardware controller to maintain the processor always under an assigned power budget. The hardware controller tries to maximize the overall performance while constraining the power consumption and maintaining a safe silicon temperature. Intel architectures implement in its CPU a hardware power controller called Running Average Power Limit (RAPL) depicted in figure 2. RAPL is a control system, which receives as input a power limit and a time window. As consequent, RAPL continuously tunes the P-states to ensure that the limit is respected in the specified time window. RAPL can scale down and up core's frequencies when the power constraint is not respected overriding the selected P-states. RAPL power budget and time window can be configured writing a Machine Specific Register (MSR) on the CPU. Maximum and minimal values for both power budget and time window are specified in a read-only architectural register. Values for both power and time used in RAPL are represented as multiple of a reference unit contained in a specific architectural register. At the machine startup, RAPL is configured using Thermal Design Power (TDP) as power budget with a 10ms time window. RAPL also provides 32bit performance counters for each power domain to monitor the energy consumption and the total throttled time. RAPL implements four power domains which can be independently configured:

- (1) *Package Domain*: this power domain limits the power consumption for the entire package of the CPU, this includes cores and uncore components.
- (2) DRAM Domain: this power domain is used to power cap the DRAM memory. It is available only for server architectures.
- (3) *PP0/Core Domain*: is used to restrict the power limit only to the cores of the CPU.
- (4) PP1/Graphic Domain: is used to power limit only the graphic component of the CPU. It is available only for client architectures due Intel server architectures do not implement graphic component into the package.

In the experimental result section, we focus our exploration on the package domain of RAPL controller because core and graphic domains are not available on our Intel architecture. DRAM domain is left for future exploration works. We also tried to modify the time windows of package domain (which can be set in a range of 1ms to 46ms in our target system) to see its impact on application performance. Our results show that this parameter does not lead to noticeable changes in the results obtained. For this reason, we report results only for the default 10ms time window configuration.

4 BENCHMARKING SCENARIO

4.1 Architecture Target

In this work, we take as architecture target a high-performance computing infrastructure, which is a Tier-1 HPC system based on an IBM NeXtScale cluster. Each node of the system is equipped with 2 Intel Haswell E5-2630 v3 CPUs, with 8 cores with 2.4 GHz nominal clock speed and 85W Thermal Design Power (TDP, [12]). As regards the software infrastructure, SMP CentOS Linux distribution version 7.0 with kernel 3.10, runs on each node of the system. We use the complete software stack of Intel systems for HPC production environment. In particular, we use Intel MPI Library 5.1 as the runtime for communication and Intel ICC/IFORT 16.0 in our toolchain. This Tier-1 supercomputer is currently classified in the Top500 supercomputer list [1]. We focus our analysis on a single node of the cluster.

4.2 Application Target

Quantum ESPRESSO (QE) [11] is an integrated suite of computer codes for electronic-structure calculations and materials modelling at the nanoscale. It is an open source package for research in molecule dynamics simulations and it is freely available to researchers around the world under the terms of the GNU General Public License. Quantum ESPRESSO is commonly used in high-end supercomputers. QE main computational kernels include dense parallel Linear Algebra (LA) and 3D parallel Fast Fourier Transform (FFT). Moreover, most of application workload is based on LA and FFT mathematical kernels which makes our exploration work relevant for many HPC codes. In our tests, we use a Car-Parrinello (CP) simulation, which prepares an initial configuration of a thermally disordered crystal of chemical element by randomly displacing the atoms from their ideal crystalline positions. This simulation consists of a number of tests that must be executed in the correct order.

4.3 Monitoring Framework

In this section, we describe in detail the monitoring framework used to profile the application and the system. Our monitoring framework is composed to two monitoring tools. The first one can monitor several hardware performance counters with a regular time stamping. The second monitoring framework is synchronized with parallel phases allowing to isolate performance and architectural metrics for each program phase. However due to the higher number of monitoring points per time unit, it can access to only on a sub-set of performance counters of the one monitored by system monitoring tool. Our monitoring frameworks have a minimal overhead, less than 1% w.r.t. application execution time.

- 4.3.1 System-aware Monitoring Tool. We use as system-aware monitoring tool Examon [2]. This monitoring tool can be used to read periodically per-core frequency, CPI and scalar/vector instructions retired. In addition, it can monitor for each socket the DRAM memory bandwidth and package power consumption using RAPL performance counters. This monitor is a simple daemon process that can access to the performance counters of the CPU using MSR read/write operations. The daemon starts at a given T_{samp} rate, in our benchmarks we use a T_{samp} of 1 second.
- 4.3.2 Application-aware Monitoring Runtime. We developed a monitor runtime to extract system information synchronously with the application flow. The runtime is a simple wrapper of the MPI library where every MPI function of each process has been enclosed by an epilogue and a prologue function. We used the MPI standard profiling interface (PMPI), which allow us to intercept all the MPI library functions without modify the application source code. The runtime is integrated in the application at linked time. Hence,

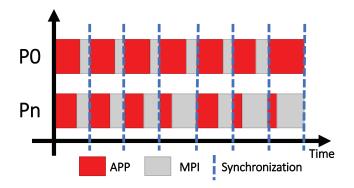


Figure 3: Phases of computation and communication identified by application-aware Monitoring Runtime

Application-aware Monitoring Runtime is able to extract information distinguishing application and MPI phases as shows in figure 3. The monitor runtime uses RDPMC and RDTSC assembly instructions to access respectively the Time Stamp Counter (TSC) and Intel Performance Monitoring Unit (PMU) counters with an overhead of few hundreds of cycles for each counter access. PMU counters are programmable through standard MSR operations which require administrative permissions and are costly in terms of access overhead. However, the counter value can be read using the RDPMC instruction directly from user space without involving syscalls. We programmed per-core PMU registers to monitor frequency, CPI, and scalar/vector instructions retired. The monitor runtime can intercept a very high number of MPI calls of the application, for this reason is not possible to use MSR operations to access low level performance counters through syscalls, which cause high-performance penalty.

5 EXPERIMENTAL RESULTS

5.1 Methodology

We run QE-CP with a configuration of 16 MPI processes with a oneto-one bind to each core of our HPC node. We start by comparing different configurations of power capping in our test environment. Initially, we split the power budget in an equal manner on both sockets, we set RAPL to maintain 48W on each socket, for a global power envelope of 96W. This test shows that the core's frequencies on different sockets are heterogeneous, suggesting that the two sockets have different inherent power efficiency. To have the same frequency among all the cores, the tested computing node needs of 11.31% higher power on socket 0. As consequence of this result, we run a set of benchmarks using the OS power manager to fix the same frequency for all the cores while monitoring the power consumption of each socket during an application run. We use these per-socket power budgets as power constraints to obtain the same frequency among all the cores. We execute again the tests using RAPL to impose these per-socket power caps and leave RAPL decides the actual frequency.

Table 1 shows the results of our set of experiments using different levels of power caps. In the first column, there are reported the target frequencies used to extract the power limits specified in the second column. Second and third columns show the sum of power consumption of both sockets using Fixed Frequency (FF) allocation and RAPL mechanisms for power capping. We can see that the power consumption is the same, so the power cap is respected and the tests are comparable. In the frequency columns are reported the average frequencies for the entire application and among all the cores. These columns show that RAPL has an average frequency

of 11.14% higher than FF but, if we take a look at the execution time (reported in next columns), FF has a lower execution time, in average 2.87% faster than RAPL. In the next sections, we will explore why FF has a lower execution time respect to RAPL which, in contrast, has a higher average frequency.

5.2 System Analysis

Figure 4 shows a time window of the system-aware monitoring tool for both power capping mechanisms while QE-CP iterates on the same computational kernel. The test reports the case of a power constraint relative to 1.5 GHz for FF and RAPL. So, the results are comparable directly.

First, we can check the correct behavior of power capping logic by looking at the core's frequencies and package power consumption (first two top plots). In the FF plot on the left part of figure 4, core's frequencies are fixed at 1.5 GHz while package power consumption floats around the average value as effect of the different application phases. In contrast, RAPL (on the right) maintains constant the power consumption for both sockets while core's frequencies changes following the current application phase. Table 1 reports a similar average power consumption for both cases, thus the power cappers are working as expected. Both benchmarks show a lower CPI when the memory bandwidth is low and SIMD instructions retired are high. In these phases, RAPL has lower frequency than the FF case as effect of the higher power demand of SIMD instructions. On the other hand, RAPL assigns higher frequencies than FF when CPI is high and this happens when the application is moving data from/to memory as proved by the high memory traffic/bandwidth reported by the Mem Ch[GB/s] plot. In these phases, the number of SIMD instructions retired are lower and, as already pointed out and shown in the RAPL plot, the core's frequencies selected by RAPL increases above average due the higher available power budget. However, increasing core's frequencies when the application is memory bound does not reflect in a consequent performance gain due the higher CPI and sub-linear dependency of application speed-up with frequency in these phases.

Hence, FF is more efficient of RAPL (shorter execution time) for two reasons: i) FF executes with higher instruction per seconds when application has high SIMD instructions density. ii) RAPL instead reduces the core's frequency in the same phase to avoid excessive power consumption. On the contrary, RAPL increases the frequency during memory bound phases obtaining a similar average power as the FF case.

5.3 Application Analysis

In this section, we take a look what happen in the system through the application-aware monitoring runtime. This runtime is able to recognize application phases marked by global synchronization points as depicted in figure 4. In the figures 5, 6, and 7 are reported the average values of performance counters of RAPL gathered into frequency operational intervals. In figure 5 is depicted the amount of time for computation (APP) and for communication (MPI) phases. Figure 6 shows the time gain for the FF respect to RAPL distributed on different frequency operational intervals. Negative values mean seconds of application time saved by FF w.r.t. RAPL. Figure 7 shows the values for CPI and SIMD instructions retired for the same phases that RAPL executes at a given frequency.

To explain the behavior that characterizes FF and RAPL, we need to look at all the three plots together. Starting from figure 6, we can recognize that in the frequency intervals 1.3 to 1.5 GHz and 1.7 to 1.9 GHz, FF obtains its highest speed up. The first gaining interval is justified by the high number of SIMD instructions retired and by the lower CPI with respect to other application phases. Indeed

	Power		Frequency			Execution Time		
	FF	RAPL	FF	RAPL	FF vs RAPL	FF	RAPL	FF vs RAPL
1.5 GHz	95.56W	94.81W	1499MHz	1766MHz	-15.11%	311.43sec	328.16sec	5.10%
1.8 GHz	111.86W	110.63W	1797 <i>MHz</i>	2144MHz	-16.22%	274.11sec	274.42sec	0.11%
2.1 GHz	122.87W	120.71W	2094 <i>MHz</i>	2323MHz	-9.86%	247.60sec	254.59sec	2.75%
2.4 GHz	134.44W	131.32W	2392MHz	2476MHz	-3.37%	231.19sec	239.65sec	3.53%

Table 1: Quantum ESPRESSO - Power Capping

2.0	Fixed Fre		RAPL					
1.8 1.8 1.7 1.6 1.5 1.5 1.5 1.4				1.9 1.8 1.7 1.6 1.5				
1.3	170	180	190	1.3	160	170	180	190
(X) 50 50 50 40 8 30 0 20			Socket 0 Socket 1	50 40 30				Socket 0 Socket 1
<u>S</u> 10	170	180	190	10	160	170	180	190
2.0 1.5 6 1.0 0.5	\mathcal{M}		\bigvee	1.5 1.0 0.5	\mathcal{N}		$\bigwedge \mathcal{F}$	
0.0 2.5 1e10 160	170	180	190 Socket 0	0.0 2.5 2.0	160	170	180	190 — Socket 0
E 1.5 OM 1.0 In 0.5			— Socket 1	1.5 1.0 0.5 0.0		$/ \setminus / \setminus$		Socket 1
0.0160	170	180	190	30.	160	170	180	190
(s) 25 (S) 25 (S			Socket 0 Socket 1	25 20 15 10 5				Socket 0 Socket 1
160	170 Time (180	190	-	160	170 Time (se	180 ec)	190

Figure 4: Time window of 50 seconds of the system monitor, every value is averaged over 1-second-interval window

from figure 7, we can notice that most of the SIMD instructions are executed with these lower frequencies by RAPL.

In the interval 1.7 to 1.9 GHz, the CPI is higher and the SIMD instructions retired are not negligible. From figure 5, we can recognize that most the application time is spent in this frequency range with a lower SIMD instructions density respect to the 1.3 to 1.5 GHz interval. Hence, high CPI, low density of SIMD instructions and high frequency suggest memory bound phases as shown by previous section. Interesting these phases run at higher frequency than the FF but leads to a performance penalty. This suggests side effects of high frequency in terms of memory contentions.

In the interval 2.0 to 2.1 GHz, RAPL has a performance gain with respect to FF. This behavior is explained by the CPI and the number of SIMD instructions retired during this phase. In this interval, RAPL has a low CPI and does not perform SIMD instructions, so this phase scales its execution with the frequency. RAPL can dynamically manages the available power budget made available by the low SIMD instructions and can increases the frequency. This leads to a consequent performance increment.

In the turbo frequency interval, RAPL performs better than FF as it is a MPI reduction phase where only the root process is active. During the reduction, all the processes except the root MPI remain in a barrier to wait the termination of the root. This is explained by the high MPI runtime time (figure 5) presents at this frequency interval. Hence, RAPL can use the power budget released by the processes in barriers to speed up the root process leading to a performance gain.

6 CONCLUSION

In this paper, we presented a novel exploration work on power capping mechanisms for power constrained HPC nodes. Differently from state-of-the-art explorations, we focused our analysis on power capping strategies used in real HPC system nodes. In details, we explored the characteristics of Intel RAPL and a fixed frequency allocation during the execution of a real scientific HPC application which is performance constrained by the power budget assigned to the node.

Our exploration explains why RAPL mechanism has an average performance penalty of 2.87% (up to 5.10%) w.r.t. the fixed frequency case, even if RAPL shows a higher average frequency for the entire application time. This proved that dynamically manage the available power budget without be aware of the application phases is not always beneficial from performance point of view. Furthermore, RAPL mechanism always increases the core's frequencies during less demanding phases to fill in the available power budget, leading to unnecessary power consumption in memory bound phases that can cause interference and as consequence performance loss.

In future works we will leverage these considerations to design a software power capping mechanisms which enforces a power cap in average on a user defined time-period. This will complement today RAPL mechanisms allowing power shifting in between application phases. For this purpose we plan to extend previously presented approaches, which combine speculative power management on short-time periods (O.S. ticks) with the ability of maintaining average power management goals on longer periods (aging periods)[15, 16].

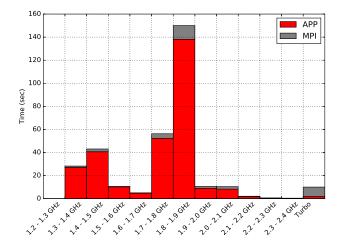


Figure 5: Sum of MPI and application time gathered at frequency operational intervals

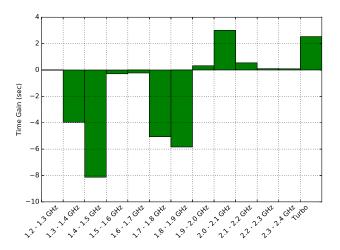


Figure 6: Time gain of FFP w.r.t RAPL gathered at frequency operational intervals

ACKNOWLEDGMENTS

Work supported by the EU FETHPC project ANTAREX (g.a. 671623), EU project ExaNoDe (g.a. 671578), and EU ERC Project MULTI-THERMAN (g.a. 291125).

REFERENCES

- $[1]\ \ 2017.\ TOP 500.Org.\ Top\ 500\ Supercomputer\ Sites.\ http://www.top 500.org.\ (2017).$
- [2] Francesco Beneventi, Andrea Bartolini, Carlo Cavazzoni, and Luca Benini. 2017. Continuous learning of HPC infrastructure models using big data analytics and inmemory processing tools. In 2017 Design, Automation & Test in Europe Conference & Exhibition (DATE). IEEE, 1038-1043.
- [3] Keren Bergman, Shekhar Borkar, Dan Campbell, William Carlson, William Dally, Monty Denneau, Paul Franzon, William Harrod, Kerry Hill, Jon Hiller, et al. 2008. Exascale computing study: Technology challenges in achieving exascale systems. Defense Advanced Research Projects Agency Information Processing Techniques Office (DARPA IPTO), Tech. Rep 15 (2008).
- [4] Andrea Borghesi, Andrea Bartolini, Michele Lombardi, Michela Milano, and Luca Benini. 2016. Predictive Modeling for Job Power Consumption in HPC Systems. In International Conference on High Performance Computing. Springer, 181–199.
- [5] Andrea Borghesi, Christian Conficoni, Michele Lombardi, and Andrea Bartolini. 2015. MS3: a Mediterranean-Stile Job Scheduler for Supercomputers-do less

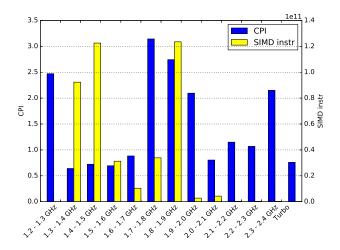


Figure 7: Average CPI and number of AVX instructions retired gathered at frequency operational intervals

- when it's too hot!. In High Performance Computing & Simulation (HPCS), 2015 International Conference on. IEEE, 88–95.
- [6] Hao Chen, Michael C Caramanis, and Ayse K Coskun. 2014. The data center as a grid load stabilizer. In Design Automation Conference (ASP-DAC), 2014 19th Asia and South Pacific. IEEE, 105–112.
- [7] Hao Chen, Michael C Caramanis, and Ayse K Coskun. 2014. Reducing the data center electricity costs through participation in smart grid programs. In Green Computing Conference (IGCC), 2014 International. IEEE, 1–10.
- [8] Jack Dongarra. 2013. Visit to the national university for defense technology changsha, china. Oak Ridge National Laboratory, Tech. Rep., June (2013).
- [9] Jonathan Eastep, Steve Sylvester, Christopher Cantalupo, Federico Ardanaz, Brad Geltz, Asma Al-Rawi, Fuat Keceli, and Kelly Livingston. 2016. Global extensible open power manager: a vehicle for HPC community collaboration toward codesigned energy management solutions. Supercomputing PMBS (2016).
- [10] Neha Gholkar, Frank Mueller, and Barry Rountree. 2016. Power tuning HPC jobs on power-constrained systems. In Proceedings of the 2016 International Conference on Parallel Architectures and Compilation. ACM, 179–191.
- [11] Paolo Giannozzi, Stefano Baroni, Nicola Bonini, Matteo Calandra, Roberto Car, Carlo Cavazzoni, Davide Ceresoli, Guido L Chiarotti, Matteo Cococcioni, Ismaila Dabo, et al. 2009. QUANTUM ESPRESSO: a modular and open-source software project for quantum simulations of materials. *Journal of physics: Condensed matter* 21, 39 (2009), 395502.
- [12] Per Hammarlund, Kajesh Kumar, Randy B Osborne, Ravi Rajwar, Ronak Singhal, Reynold D'Sa, Robert Chappell, Shiv Kaushik, Srinivas Chennupaty, Stephan Jourdan, et al. 2014. Haswell: The fourth-generation intel core processor. *IEEE Micro* 34, 2 (2014), 6–20.
- [13] Emma Jane Hogbin. 2015. ACPI: Advanced Configuration and Power Interface. (2015).
- [14] Aniruddha Marathe, Peter E Bailey, David K Lowenthal, Barry Rountree, Martin Schulz, and Bronis R de Supinski. 2015. A run-time system for power-constrained HPC applications. In International Conference on High Performance Computing. Springer, 394–408.
- [15] Pietro Mercati, Andrea Bartolini, Francesco Paterna, Tajana Simunic Rosing, and Luca Benini. 2014. A linux-governor based dynamic reliability manager for android mobile devices. In Design, Automation and Test in Europe Conference and Exhibition (DATE), 2014. IEEE, 1–4.
- [16] Pietro Mercati, Francesco Paterna, Andrea Bartolini, Luca Benini, and Tajana Šimunić Rosing. 2017. WARM: Workload-Aware Reliability Management in Linux/Android. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems 36, 9 (2017), 1557–1570.
- [17] Venkatesh Pallipadi and Alexey Starikovskiy. 2006. The ondemand governor. In Proceedings of the Linux Symposium, Vol. 2. sn, 215–230.
- [18] Osman Sarood, Akhil Langer, Abhishek Gupta, and Laxmikant Kale. 2014. Maximizing throughput of overprovisioned hpc data centers under a strict power budget. In Proceedings of the International Conference for High Performance Computing Nature Plans and Analysis IEEE Pages 2017. 2819.
- puting, Networking, Storage and Analysis. IEEÉ Press, 807–818.
 [19] Alina Sirbu and Ozalp Babaoglu. 2016. Predicting system-level power for a hybrid supercomputer. In High Performance Computing & Simulation (HPCS), 2016 International Conference on. IEEE, 826–833.
- [20] Huazhe Zhang and H Hoffman. 2015. A Quantitative Evaluation of the RAPL Power Control System. Feedback Computing (2015).